

# **EFMS Kafka Integration Architecture Documentation**

<b>1) Overview</b>	<b>3</b>
<b>2) Architecture Goals and Design Principles</b>	<b>3</b>
<b>3) System Context (What exists in EFMS)</b>	<b>4</b>
System Context Diagram	5
<b>4) Kafka Infrastructure (Docker Cluster)</b>	<b>5</b>
4.1 Cluster Type	5
4.2 Listeners (Why there are two ports per broker)	5
4.3 SASL/PLAIN authentication	6
Kafka Deployment Diagram	7
<b>5) Kafka in EFMS Backend (How your code uses Kafka)</b>	<b>7</b>
5.1 Configuration (in-memory, currently)	7
5.2 Producers (publish workflow events)	8
5.3 Consumers (background workers)	8
Consumer Groups Diagram	9
<b>6) SignalR in EFMS (Real-time UI control)</b>	<b>9</b>
6.1 EFMSHub responsibilities	9
6.2 UserPresenceTracker	10
<b>7) Key Workflows (End-to-end)</b>	<b>10</b>
7.1 OTP Generation + Email Delivery (DB + Kafka)	10
Sequence Diagram	11
7.2 Signup Approval/Rejection Email (Kafka)	12
Sequence Diagram	13
7.3 User Deactivation → Force Logout (SignalR) + Backend Enforcement (Kafka)	14
Sequence Diagram	15
<b>8) Deployment Notes (Docker Hosting Reality)</b>	<b>16</b>
8.1 If EFMS backend runs on the HOST machine (current config)	16
8.2 If EFMS backend is moved into Docker (important when “host using docker”).	16
Host vs Docker Connectivity Diagram	17
<b>9) Reliability &amp; Operational Behavior</b>	<b>18</b>
9.1 Message durability	18
9.2 Re-processing	18
9.3 Offsets and duplicates (what to expect)	18
<b>10) Observability (How you troubleshoot)</b>	<b>19</b>
10.1 Kafdrop	19
10.2 EFMS logging	19
<b>11) Event Contracts (What EFMS publishes)</b>	<b>20</b>
11.1 OTP event (current shape)	20
11.2 SignupDecisionEvent (your class)	20
11.3 Session expiry message (current)	20
<b>12) Why this Architecture Works Well for EFMS</b>	<b>21</b>

## **1)Overview**

**This document describes how EFMS uses event streaming and real-time messaging to support background workflows and immediate user updates.**

- Kafka is used as a reliable event channel between EFMS components. Instead of doing slow work inside API requests (such as sending email), EFMS publishes an event and processes it in the background.
- SignalR is used to send immediate updates to the user interface (example: forcing a user logout across all active sessions).

**This design supports:**

- Faster API responses (work is handled asynchronously)
- Reliable processing (events can be re-processed if required)
- Better user experience (real-time notifications and session control)

## **2)Architecture Goals and Design Principles**

**Goals:**

- Responsiveness: Requests shouldn't wait for email/SMS delivery.
- Reliability: Workflow signals must survive service restarts and temporary failures.
- Scalability: Increase throughput by scaling consumers (instances / consumer groups).
- Traceability: Inspect topics and messages via Kafdrop for debugging/audit support.
- Immediate UI control: Critical actions must reflect instantly on connected clients.

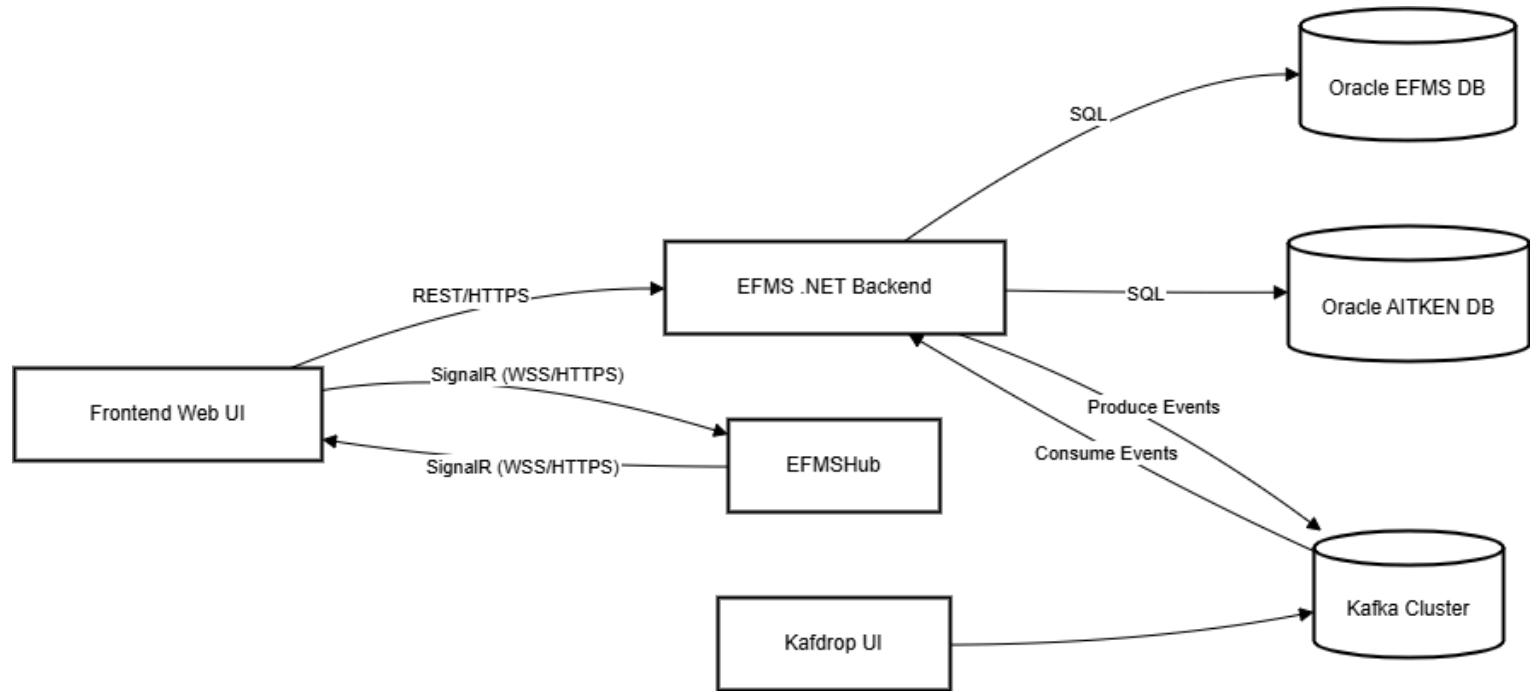
### **Design Principles:**

- Oracle DB is the system of record (OTP tokens, user status, approvals).
- Kafka carries workflow signals, not data ownership.
- SignalR = “update UI now”, Kafka = “process consistently in background.”

### **3)System Context (What exists in EFMS)**

- Frontend (Web UI): Calls EFMS APIs and maintains a SignalR connection.
- Backend (.NET API): Business logic + DB access + Kafka producers/consumers + SignalR hub.
- Oracle Databases:
  - EFMS (auth, OTP, users, approvals, setup)
  - AITKEN (fund-related domain data)
- Kafka Cluster: durable event pipeline for background workflows.
- Kafdrop: operational visibility (topics, partitions, messages).

## System Context Diagram



**Image 1**

## 4) Kafka Infrastructure (Docker Cluster)

### 4.1 Cluster Type

Your docker-compose creates a 3-broker Kafka cluster using KRaft (brokers are also controllers). This improves simplicity (no ZooKeeper) and supports replication and fault tolerance.

### 4.2 Listeners (Why there are two ports per broker)

**You configured dual listeners to support both:**

- Host machine apps (your EFMS backend running on localhost)
- Docker network apps (kafdrop + any containerized services)

### **Per broker:**

- **SASL\_PLAINTEXT://localhost:9092/9093/9094** → used by EFMS backend running on host
- **PLAINTEXT://kafkaX:19092/19093/19094** → used by Docker services (Kafdrop)
- **CONTROLLER://...** → internal KRaft controller quorum traffic

## **4.3 SASL/PLAIN authentication**

### **Your server\_jaas.conf defines Kafka users:**

- producer / prod-secret
- consumer / cons-secret
- broker internal user

### **Efms uses:**

```
SecurityProtocol = SASL_PLAINTEXT
SaslMechanism = PLAIN
SaslUsername = KafkaAuth.Username // "producer"
SaslPassword = KafkaAuth.Password // "prod-secret"
```

## Kafka Deployment Diagram

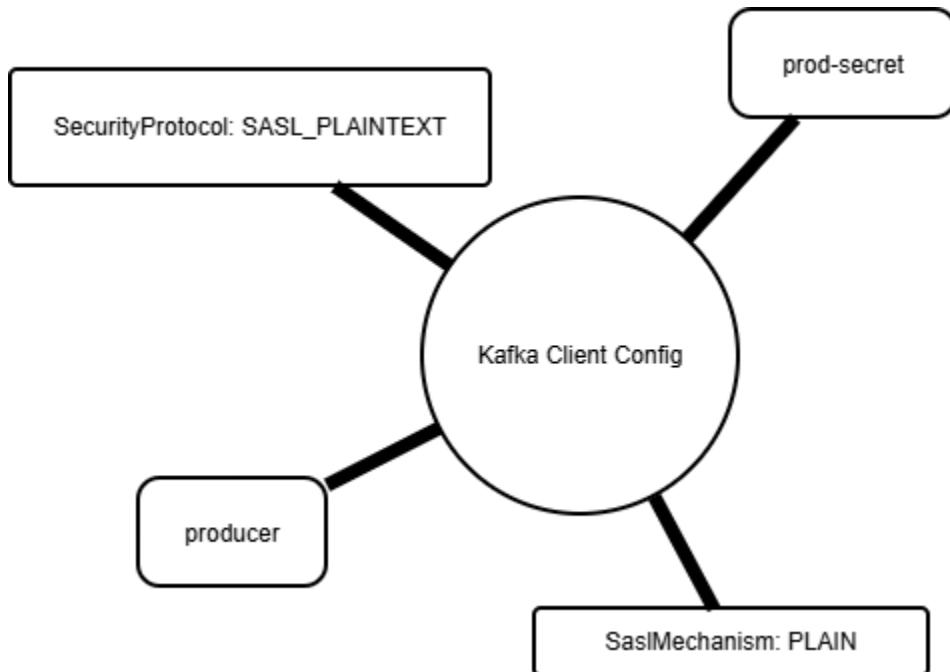


Image 2

## 5) Kafka in EFMS Backend (How your code uses Kafka)

### 5.1 Configuration (in-memory, currently)

Your backend config sets:

BootstrapServers = localhost:9092,localhost:9093,localhost:9094

GroupId = efms-session-expiry-group (with suffixes per consumer)

Topics:

- otp-service
- signup-rejection
- session-expiry

## 5.2 Producers (publish workflow events)

**In the Code :**

- **KafkaProducerService** → publishes OTP events to otp-service
- **SignupDecisionProducer** → publishes approval/rejection decision to signup-rejection
- **SessionExpiryProducer** → publishes logout/session-expiry signal to session-expiry

## 5.3 Consumers (background workers)

**Hosted services in your backend:**

- **EmailOtpConsumer** (consumer group: ...-email)  
Reads otp-service → sends email OTP
- **SmsOtpConsumer** (consumer group: ...-sms)  
Reads otp-service → placeholder for SMS provider integration
- **EmailSignupDecisionConsumer** (consumer group: ...-signup)  
Reads signup-rejection → sends approval/rejection email
- **SessionExpiryWorker** (group: efms-session-expiry-group)  
Reads session-expiry → triggers ExpiryPushService.Schedule(userId, expiresAt)

## Consumer Groups Diagram

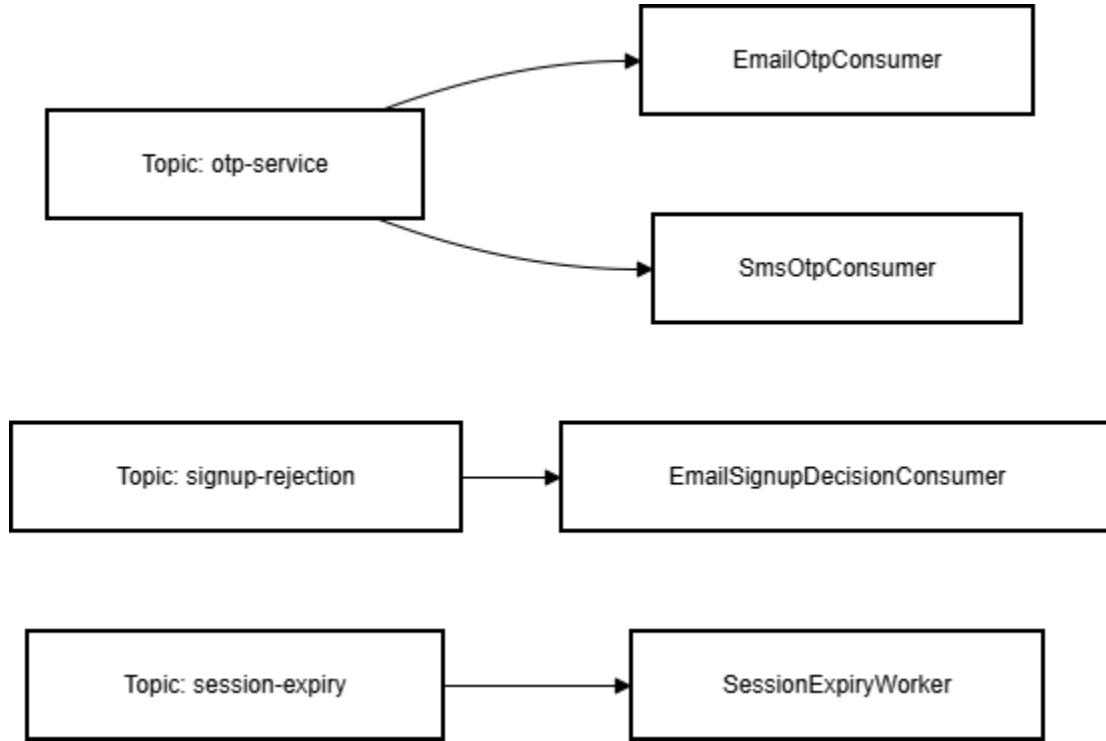


Image 3

## 6) SignalR in EFMS (Real-time UI control)

### 6.1 EFMSHub responsibilities

#### **On connect:**

- identifies user via JWT (token from query string or cookie)
- adds connection to a company group: company\_{companyId}
- tracks connections via UserPresenceTracker

#### **On disconnect:**

- removes connection from tracker

## **6.2 UserPresenceTracker**

**This is how EFMS knows all active connections for a user so it can:**

- send targeted events (only that user)
- force logout across all tabs/devices immediately

# **7) Key Workflows (End-to-end)**

## **7.1 OTP Generation + Email Delivery (DB + Kafka)**

### **What happens**

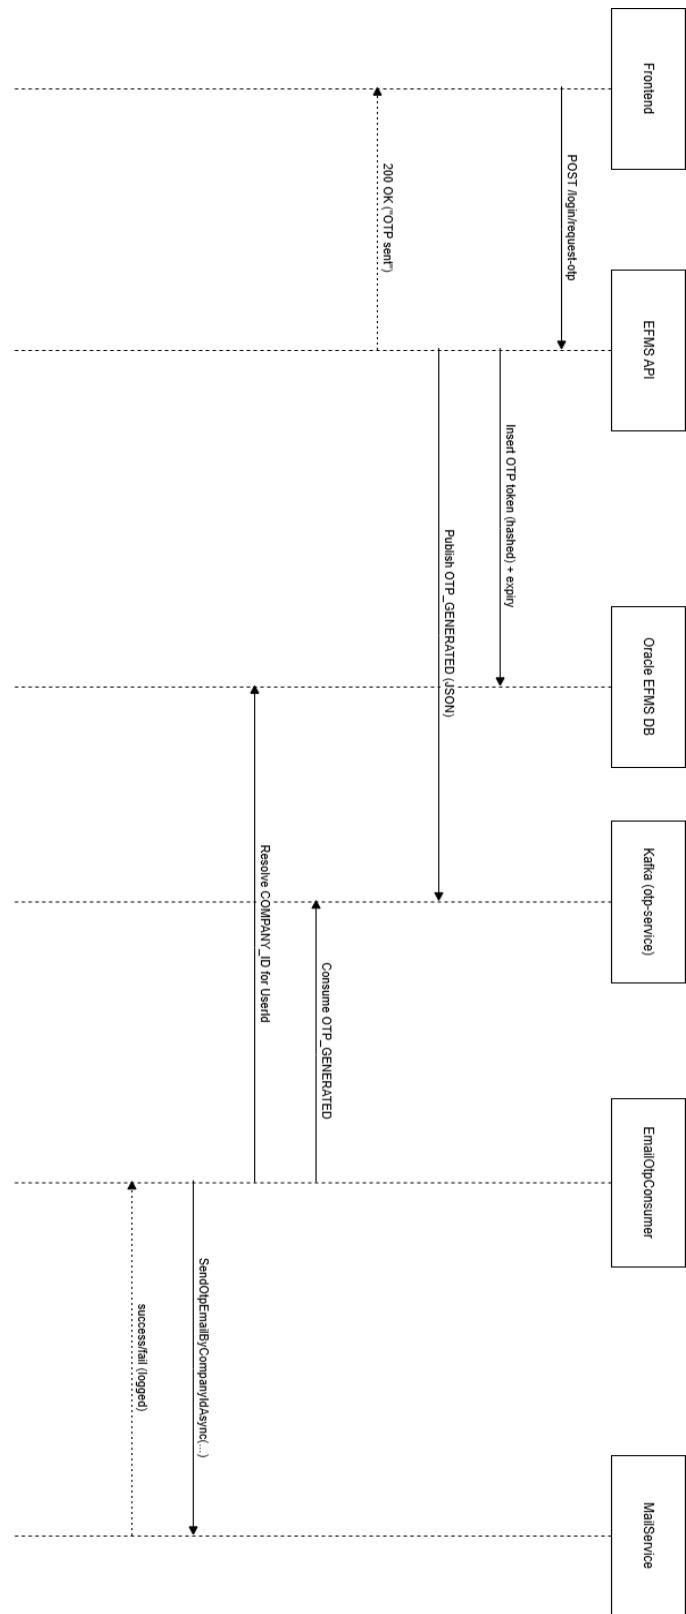
- API generates OTP and stores hashed OTP in EFMS DB.
- API publishes an OTP\_GENERATED event to Kafka (otp-service).
- EmailOtpConsumer consumes it, resolves companyId (DB lookup), sends OTP email.

### **Why this design**

- API returns fast (“OTP sent”) without waiting for email.
- Email sending is retryable and inspectable.

# Sequence Diagram

Image 4



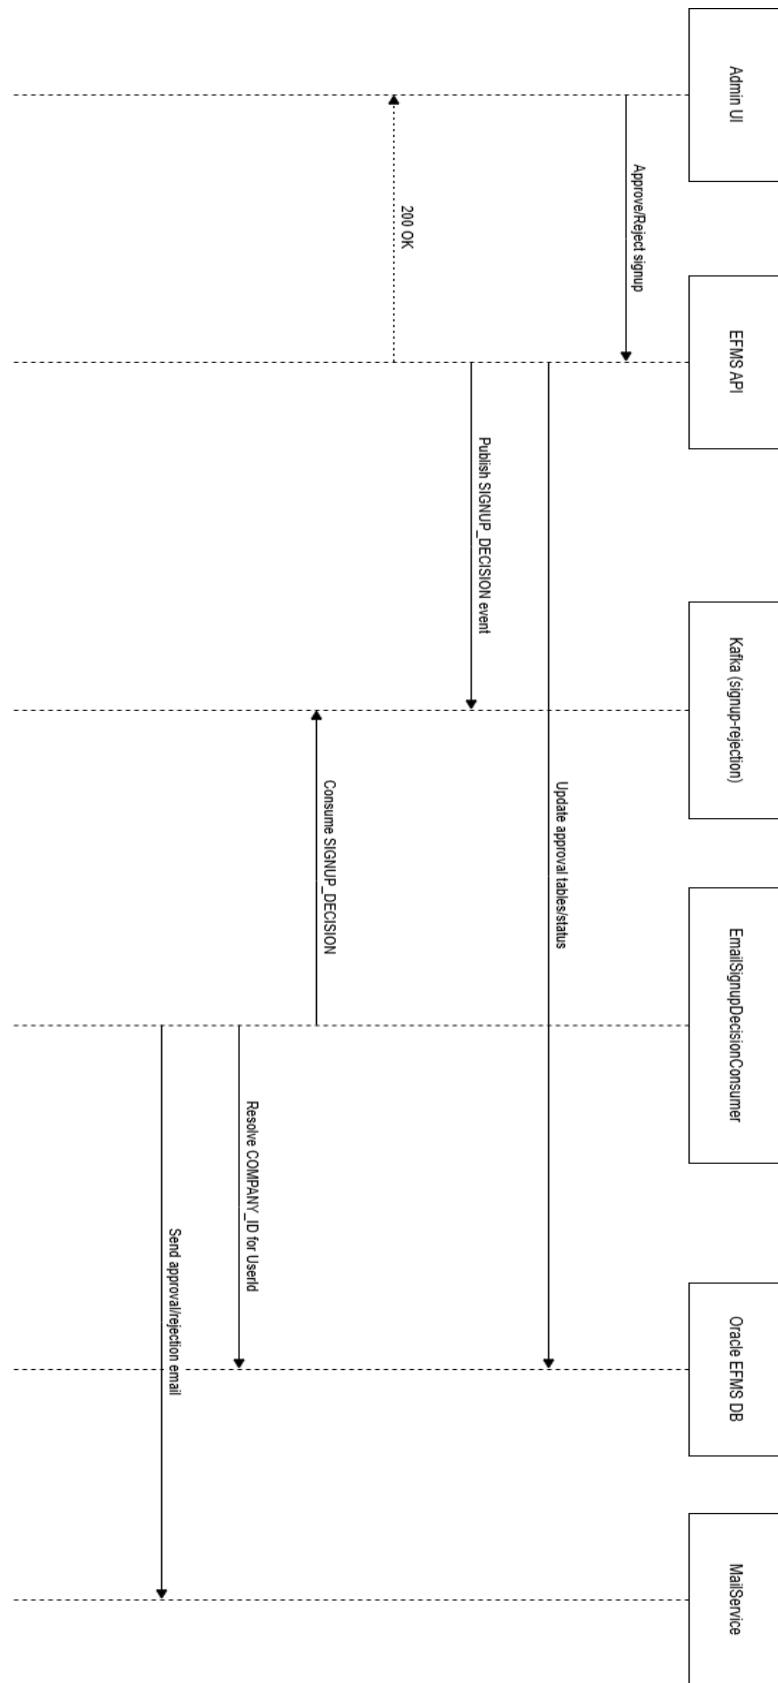
## **7.2 Signup Approval/Rejection Email (Kafka)**

### **What happens**

- After approval flow, EFMS publishes SignupDecisionEvent to topic signup-rejection ( it currently carries both approved and rejected decisions).
- EmailSignupDecisionConsumer sends the appropriate email template.

# Sequence Diagram

Image 5



## **7.3 User Deactivation → Force Logout (SignalR) + Backend Enforcement (Kafka)**

### **What happens**

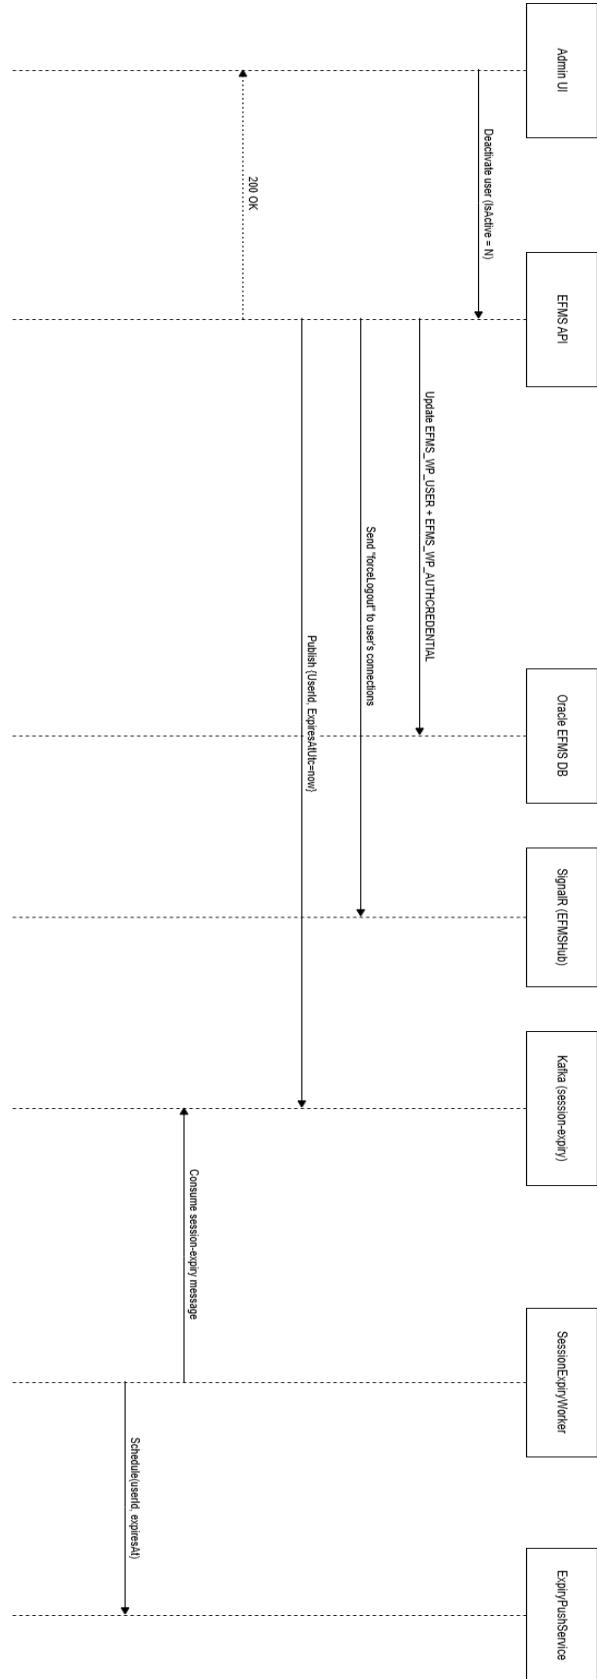
- Admin/API deactivates user in DB (EFMS\_WP\_USER + EFMS\_WP\_AUTHCREDENTIAL).
- EFMS sends SignalR forceLogout to all active connections of that user (immediate UI action).
- EFMS publishes Kafka session-expiry event (consistent backend enforcement).
- SessionExpiryWorker receives it and triggers expiry scheduling through ExpiryPushService.

### **Why both are used**

- SignalR handles the “close user session now” experience.
- Kafka ensures “even if SignalR misses a client” the backend still processes the expiry event reliably.

# Sequence Diagram

Image 6



## **8) Deployment Notes (Docker Hosting Reality)**

### **8.1 If EFMS backend runs on the HOST machine (current config)**

**Use:**

- Kafka:BootstrapServers = localhost:9092,localhost:9093,localhost:9094

Because those are SASL\_PLAINTEXT ports exposed to host

### **8.2 If EFMS backend is moved into Docker (important when “host using docker”).**

**Inside Docker, localhost means the container itself, not the host.**

**So switch bootstrap servers to the docker listener addresses:**

- kafka1:19092,kafka2:19093,kafka3:19094 (PLAINTEXT)

**Recommended approach**

- For container-to-container communication, typically enable SASL on the docker listener too (or keep PLAINTEXT only for dev).
- If keep the docker listener as PLAINTEXT (as now), ensure this is restricted to internal networks only.

## Host vs Docker Connectivity Diagram

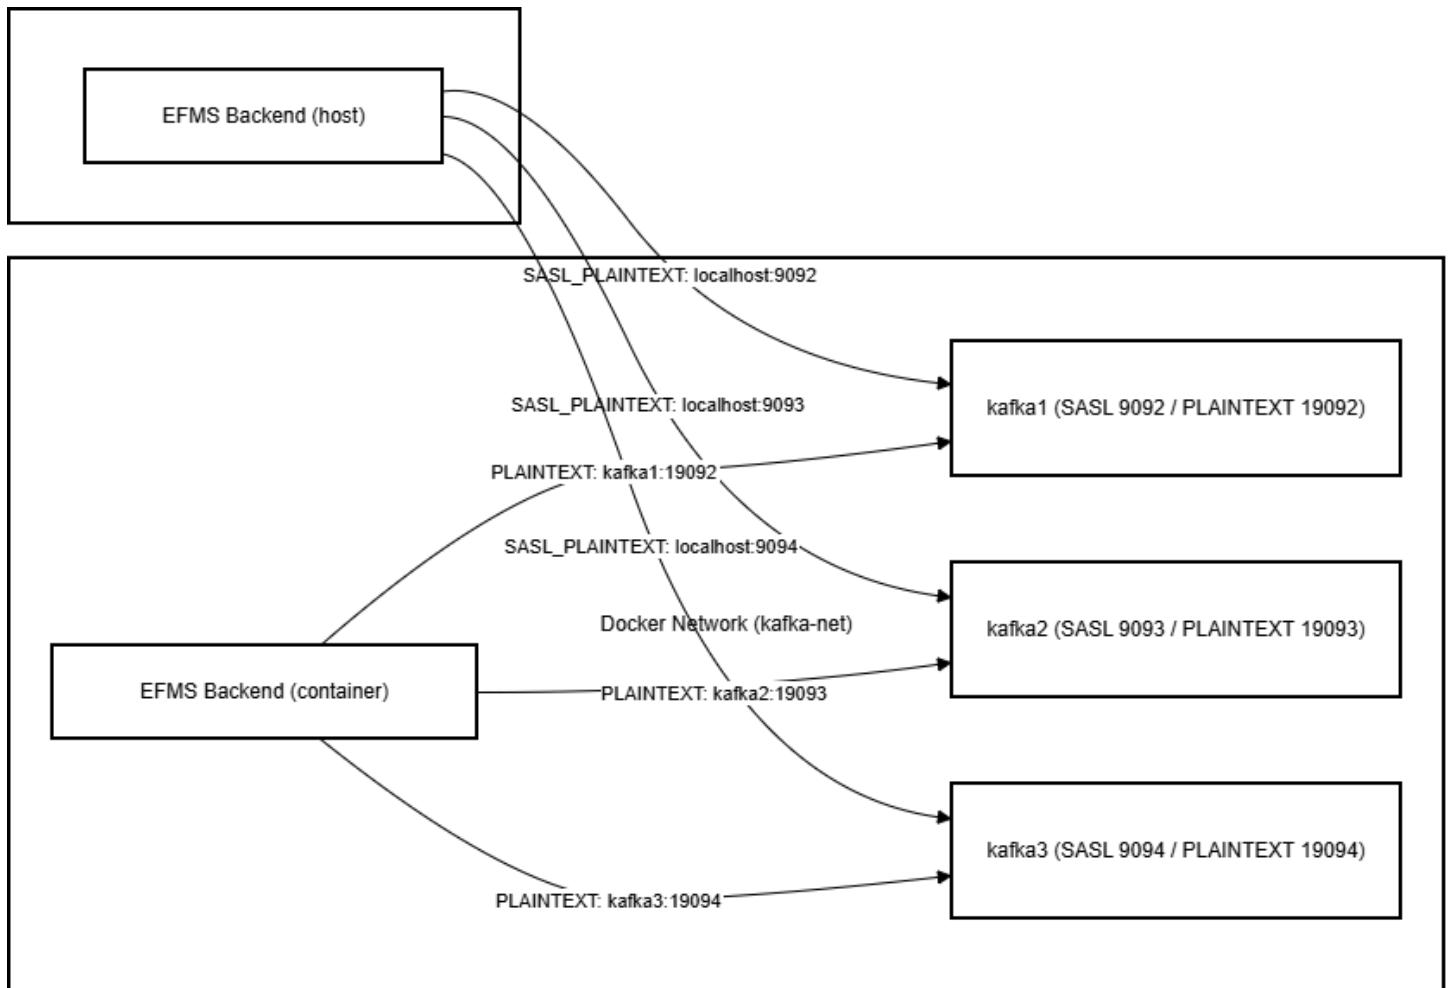


Image 7

## **9) Reliability & Operational Behavior**

### **9.1 Message durability**

Kafka stores messages on disk (replicated across brokers in your setup). This means:

- If EFMS restarts, consumers continue from last committed offset.
- If email service is temporarily down, messages remain in topic.

### **9.2 Re-processing**

Because events remain in Kafka, you can re-process by:

- using a new consumer group id (starts fresh depending on offsets policy)
- or resetting offsets (admin operation)

### **9.3 Offsets and duplicates (what to expect)**

Kafka consumers may deliver messages more than once in failure scenarios. That means:

- Email sending should be treated as **at-least-once** delivery unless you add idempotency controls.
- For OTP: duplicates are not fatal, but can confuse users. A common improvement is to include an **OtpId** in the event and store “sent status” in DB.

## **10) Observability (How you troubleshoot)**

### **10.1 Kafdrop**

**Kafdrop helps you:**

- see topic list
- confirm partitions/replication
- inspect messages (JSON payloads)
- verify consumer groups and lag

### **10.2 EFMS logging**

**Your hosted services log:**

- “Listening...” startup lines
- JSON parse errors
- send failures
- consume errors

**Recommended minimum logs per event**

- event type
- user id
- topic + partition + offset
- timestamp
- outcome (success/failure + reason)

## **11) Event Contracts (What EFMS publishes)**

### **11.1 OTP event (current shape)**

**EFMS publishes a JSON payload like:**

- Event, Context, UserId, CompanyId, Channel, Target, Otp, ExpiresAt
- Improvement suggestion (still simple):  
Add:
  - EventId (GUID)
- OtpId (DB key)  
So consumers can deduplicate.

### **11.2 SignupDecisionEvent (your class)**

**Fields:**

- Event = SIGNUP\_DECISION
- UserId, Email, Decision, Response, ApprovedBy, ApprovedName, Timestamp

### **11.3 Session expiry message (current)**

**Fields:**

- UserId
- ExpiresAt (UTC)

## **12) Why this Architecture Works Well for EFMS**

- OTPService stays fast: it writes OTP to DB and emits an event. Email sending becomes a background concern.
- Approval messaging is consistent: decision emails don't slow down admin actions.
- Deactivation is immediate AND reliable:
  - SignalR forces logout for connected users instantly.
  - Kafka ensures backend enforcement even if SignalR misses a client.